

## COMP4107 - Assignment 3

Student Name: Yunkai Wang  
Student Number: 100968473

Student Name: Jules Kuehn  
Student Number: 100661464

Fall 2018

1. a) See q1.py for code. To see a comparison of input and output to the network, uncomment `printVector()` statements.

In general, the fewer training images were used, the better the Hopfield network performed. These results suggests that a Hopfield network is a very poor choice for classification.

A major factor in the accuracy was the selection of training images. Since fewer training images performed better, a random selection of a representative 1 or 5 leaves much to chance. To address this, we selected the training images with k-means, taking the images closest to the 2 cluster centers. This results in a selection of training images in which all 1's are very close to the center 1, and all 5's are very close to the center 5. This worked better than random selection, and didn't show the same sharp decline in classification accuracy as more training images were added.

Alternately, we tried taking as many cluster centers as training images, finding the closest image for each cluster center. This should result in selecting training images that maximally differ from each other. This worked quite poorly as the quite different images caused further degeneration.

The graphs below compare classification accuracy against the number of images used to train the network (between 1 and 5 of each digit). Each graph shows the average classification accuracy over 5 trials, with 50 test images.

- b) The code to implement the Storkey improvements are also included in our q1.py code.

As shown below, the Storkey algorithm does not substantially improve the results for the classification, even for the best case shown above (where)

- c) We can further show that 8 neurons is a good choice by noting that early stopping has no effect. Our experiments have not been able to demonstrate an overtraining phenomenon with even 50 hidden neurons, but a premature early stop can be seen in Figure 6, and the table below.

Figure 1: Slow convergence

Figure 2: Premature early stop

The network with 50 hidden neurons also is very slow to converge to the target MSE - even using the RMSPropOptimizer - compared to when using 8 hidden neurons (Figure 5).

Figure 3: Reaching training goal

Figure 4: Early stopping not triggered

Figure 7 shows the MSE at each epoch near where the training goal is reached. This captures the expected trend - that the training error will decrease consistently, while the test error (9x9 grid) has more variation. The validation data is most different from the training data (as it is randomly sampled from the function), so it is expected that its error would also have the highest variance.

Figure 8 shows that having an early stopping mechanism for this function makes no difference, as the early stopping mechanism is never triggered (8 hidden neurons, 4000 epochs, 0.02 LR). Thus, the two contours are virtually identical, differing only by the stochastic elements of the experiments (initial weights, order of training data).

## 2. Instructions on how to run q2.py:

Running q2.py without changing anything will run all three parts. This may take a while as it runs 10 experiments for each number of hidden neurons, and after one is finished, a new neural net is created to run the second one, etc. Therefore, it will take about 1 min for it to finish and plot the graphs. To run a single part only, comment out the other lines in the `experiment()` function.

We started by using the default learning rate for AdamOptimizer (0.001). This took 400-700 epochs to converge to the result. However, if we use 0.01 as the learning rate, it converges faster – within 100 epochs every time. Therefore, we chose 0.01 as the learning rate for this question.

- (a) We took the averages of 10 experiments for each of 5, 10, 15, 20, and 25 hidden neurons. The results are shown below.

Figure 5: Percentage of Recognition Errors

The table below corresponds to the experiment shown in the graph above.

- (b) 15 is chosen as the number of hidden neurons, since 15 is more accurate than 10, but using more than 15 neurons is not much better (if at all). When we used 15 neurons in the hidden layer, the number of epochs to reach zero errors is much less than is shown in the assignment question, which we think is fine. Perhaps this is due to a different learning rate.

Figure 6: Training on perfect data

Figure 7: Re-training after noisy data

- (c) The network trained with both perfect and noisy data performed just as well on perfect data, and better on noisy data. The chart below is for the network with 15 neurons.

Figure 8: Percentage of Recognition Errors

### 3. Question 3

- (a) For part a, here are the parameters and dimension of the SOM computed. We used only 1024 samples to train the SOM, the reason for doing this is just to shorten the running time. Otherwise the SOM takes too long to converge. Since we are using 1024 samples, so we used a SOM of size  $12 \times 12$ . The  $\sigma$  value is 1, and the learning rate is 0.1. Here are the graphs generated by using these parameters:

It's clear that as time proceed, the SOM starts to group all the 1's together, and all

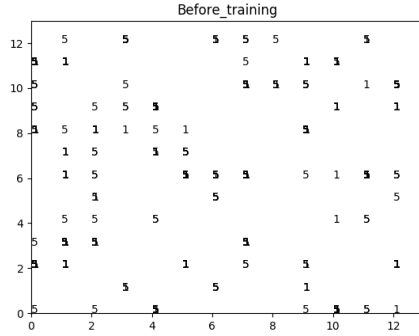


Figure 9: Training on perfect data

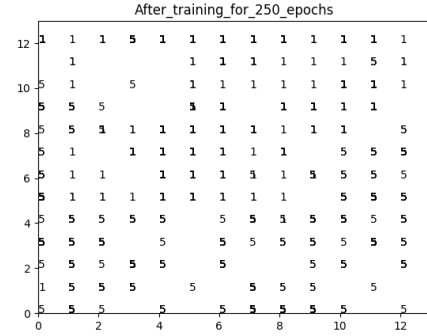


Figure 10: Re-training after noisy data




Figure 11: Re-training after noisy data

the 5's together, where initially the SOM was initialized randomly.

- (b) Here are the plots for the K-means solutions with different number of clusters.

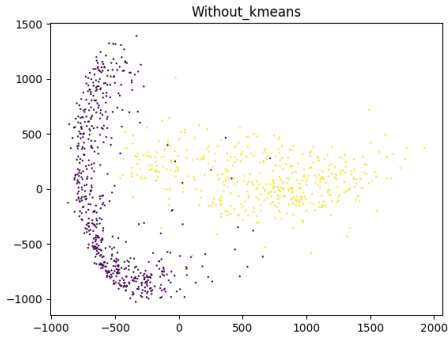


Figure 12: Training on perfect data

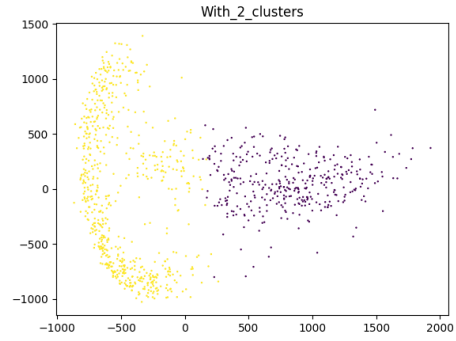


Figure 13: Re-training after noisy data



Figure 14: Re-training after noisy data

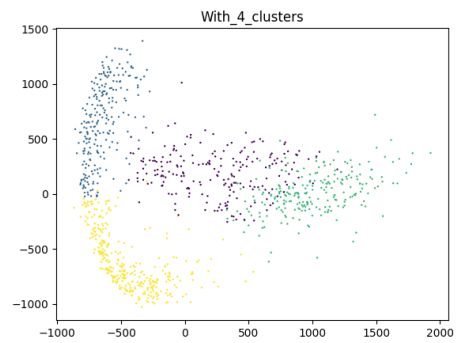


Figure 15: Re-training after noisy data



Figure 16: Re-training after noisy data

#### 4. Question 4

- (a) We tried this question with two hidden layers, and we tried a different number of hidden layers for both layers. But the accuracy doesn't change for the different number of layers that we tried. The accuracy is always 9% for all values that we tried. There may be some problem with the implementation we had.

```
>>> print(experiment(h1_size=200, h2_size=50))
0.09003436426116838
>>> print(experiment(h1_size=200, h2_size=50, learning_rate=0.01))
0.09003436426116838
>>> print(experiment(pca=True, h1_size=200, h2_size=50, learning_rate=0.01))
0.09003436426116838
```

Figure 17: Training on perfect data