

# Judging a book by its title

Yunkai Wang

Carleton University, School of Computer Science

Jules Kuehn

Carleton University, School of Computer Science

Recently Convolutional Neural Networks (CNN) have been well-studied. It has been shown that they can achieve incredible results on tasks like sentence classification (Kim, 2014; Zhang and Wallace, 2015). Also using CNN, one study tried to categorize a book based on its cover picture (Iwana et al., 2016). "However, classification of books based on the cover image is a difficult task." (Iwana et al., 2016, p. 6). While the relationship between the book's cover image and category does exist, it is difficult to classify using a CNN as many books have misleading cover images. Based on the idea of categorizing books, and using the same dataset, we experiment to determine if we can use a CNN to instead learn the potential relationship between a book's title and its categories. For this project, we will compare how accurately this task can be performed with different model configurations such as MLP and CNN (Britz, 2015), using matrices that represent the titles, where each row is a word embeddings (Mikolov et al., 2013; Karani, 2018; Britz, 2015). We will also evaluate training word embeddings from the dataset, compared with using pre-trained embeddings.

## I. Introduction

Book titles give the reader a first impression of what the book may be about, and most of the time, a good book title will attract more readers to buy and read the book. A title effectively differentiates the book from others (Peterson, 2018). But what does the book title really tell you? What can you say about the book by simply looking at the book title? If a title contains the word 'calendar', then most people will agree that it is a calendar. But what if the book title doesn't contain a specific word that reveals its category? If you are given the book title "The Three-Body Problem" with no prior knowledge, how will you categorize the book? Will the book title be sufficient for a neural network that has been trained with titles and categories to categorize that book as a science fiction, or will the network determine that it is a physics book? Frankly, we didn't know the answer to this question. Here, we conduct this experiment to see if our network is able to learn the potential relationship between a book's title and its category - and what model is best suited to the task.

## II. Background/Related work

### Dataset

*Book32* is a dataset of book information scraped from Amazon. It was published by the researchers who conducted the experiment on testing the relationship between a book's cover image and its category (Iwana et al., 2016; uchidalab, 2018). The set consists of detailed information of 207572 books from 32 different categories. Some of the books are easy to classify, like 'calendar' and 'law', which are not similar to the other categories, or they contain some specific

keyword among most books from the same category. However, many categories are similar, like 'Christian Books and Bibles' and 'Religion and Spirituality.' It is these categories that are the difficult part for the CNN to correctly classify. The dataset is not automatically splitted into training set and test set, so we use *sklearn* library to split the data.

### Loading and transforming the data

The data was supplied as a comma separated values file, with more fields than we required, such as product page URLs, product number, and author.

1. First we loaded all the original data line by line, and extracted each book's title and its category, as that was the only information we fed to our neural network.

2. Most titles are of different lengths, and the maximum title length is 96. We could have expanded the titles so that they all have length 96. However, as Figure 1 shows, most of titles have length less than 26. Therefore, we choose the maximum length to be 26 instead of 96. We used *pad\_sequence* function from *keras* library to accomplish the task, which simply appends 0 to the end of the titles that are shorter than the given length, and truncate long titles after 26 words.

3. With the help of *train\_test\_split* function, we split the training data and validation data. We used 10% of the original dataset for validation data.

4. We converted the titles into matrices that can represent these titles. More details about this follow in "Representing titles using matrices".




Figure 1. Distributions of title lengths

## Libraries

We used the following libraries to avoid reimplementing a lot of code which is not the subject of this experiment:

- Keras: We used keras library primarily, including the data preprocessing functions and the functions that create the model.
- Scikit-learn: We used sklearn to split the training and testing dataset, and the TSNE function to project the category means to a 2D plot.
- Tensorflow: Tensorflow was originally used in `mlp.py` and `cnn.py` files, and we used the models defined in those files to see how they perform on this task. It is also the backend for Keras.
- Numpy was used for various matrix processing tasks.
- Matplotlib was used to create all the graphics in this report.

## Representing titles using matrices

In order to apply CNNs to natural language processing (NLP) related tasks, the input text must be represented as matrices of numbers (Britz, 2015). We can use the two following mechanisms to represent the words in the titles:

- One-hot: Unique words are assigned indices in a lookup table. Each word in each title is then encoded into a one-hot vector, having the same dimensionality as the vocabulary size. This can be done using the `one_hot` function that is available in *keras* library.
- Word embeddings: Using a system such as *word2vec* ("Word2vec", 2018) or *GloVe* ("GloVe:

Global Vectors for Word Representation", 2014) we can generate low-rank approximations of each unique word with semantically meaningful relationships in their d-dimensional vector space (Karani, 2018).

With each word as a vector, a title can be represented as a matrix in which every row is a word vector. The first row represents the first word in the title, and so on.

## III. Problem Statement

We want computers to be able to learn the patterns in book titles, such that in the future, the computer will be able to predict the category of a new book given only its title (with good accuracy). It is almost impossible to manually design a rule-based program that will handle the categorization problem. The titles and the categories don't have an obvious relationship which we can use to linearly separate the titles into categories; there may be infinitely many potential relationships. We expect that using CNN to learn the underlying relationship between the titles and categories will achieve a high accuracy in the categorization task.

Specifically, we want to determine empirically if the CNN is better suited for this task than the MLP model, and whether using pre-trained word embeddings will help to achieve higher accuracy.

## IV. Model

We began with two naive implementations, in which each "word vector" was an arbitrary unique integer. Each title is then a vector of these integers (one for each word), zero padded to the length of the longest title (96 words). This representation of the titles has little semantic value as the integers representing the words are arbitrary.

The first naive model was a fully connected MLP with the following configuration. Output size of each layer is indicated in parentheses.

- Input (96)
- FC Sigmoid (625)
- FC Sigmoid (300)
- FC Softmax (32)

The second naive model was based on a good model for the CIFAR-10 classification.

- Input (96,1)
- Conv ReLU: 3 kernel, 1 stride (96,32)
- Conv ReLU: 3 kernel, 1 stride (96,32)

- Max pool: 2 kernel, 2 stride (48,32)
- Dropout: pKeep 0.8 (48,32)
- Conv ReLU: 3 kernel, 1 stride (48,64)
- Max pool: 2 kernel, 2 stride (24,64)
- Conv ReLU: 3 kernel, 1 stride (24,64)
- Max pool: 24 kernel, 24 stride (1,64)
- Dropout: pKeep 0.8 (1,64)
- Output Softmax (32)

A better word representation is to embed each word in a  $n$ -dimensional space, where the position of each word reflects its meaning in relation to the other words. Each word is then represented by a dense vector of length  $n$ , with similar words having similar vectors. Similarity between words is inferred from context. Many pre-trained embeddings are available, trained through Word2Vec or GloVe on datasets pulled from Wikipedia, Twitter, or other sources. An embedding can also be trained from scratch, or initialized to a pre-trained dataset then trained further. We compare these three approaches.

The first embedding model trained the embeddings from scratch, using only the book titles from the training set. Note that we truncated the titles from a maximum of 96 words to 26 words. We chose 32 for the embedding dimension, based on the rule of the thumb that the embedding should be the fourth root of the vocabulary size. There were roughly 70,000 words in the training dataset, which suggest that the embedding dimension should be 16, but we found that 32 offered a slight improvement in practice.

- Input (26, 75094)
- Embedding (26, 32)
- Flatten (832)
- Dropout 0.5 (832)
- Output Softmax (32)

The second embedding model loaded 400,000 pre-trained 100-dimensional word-vectors from the GloVe.6B dataset. These word vectors were trained on Wikipedia and Gigaword, a newswire dataset. The model is identical to EmbedTrain except that the embeddings are pre-trained and fixed, with the embedding dimension increased to 100. Dropout was also removed for this model.

- Input (26, 75094)
- Embedding (26, 100)

- Flatten (2600)
- Output Softmax (32)

The third embedding model is identical to EmbedTrain except that the embeddings are initialized to the GloVe dataset as seen in EmbedGloveFixed. Embeddings are then retrained on the book titles.

For each of the three embedding models, we also tested 2 variations to compare the performance of MLP vs CNN on the title embeddings. The first variation adds two fully connected (FC) layers:

- Input (26, 75094)
- Embedding (26, 32)
- Flatten (832)
- Dropout 0.5 (832)
- FC ReLU (1024)
- Dropout 0.5 (1024)
- FC ReLU (512)
- Dropout 0.5 (512)
- Output Softmax (32)

The second variation adds a single convolutional layer, followed by maxpool over the entire length and finally a single FC layer:

- Input (26, 75094)
- Embedding (26, 32)
- Dropout 0.5 (26, 32)
- **FC ReLU: 3 kernel, 1 stride (24, 512)**
- Max Pooling: 24 kernel (512)
- Dropout 0.25 (512)
- FC ReLU (512)
- Dropout 0.5 (832)
- Output Softmax (32)

The GloVe pre-trained (and fixed embedding) model variations are the same as above, except without the first dropout layer. The GloVe trainable models include the first dropout layer. All GloVe models necessarily have an embedding length of 100.

#### IV. Implementation

The naive models were implemented using the course-provided `cnn.py` and `mlp.py` code. All other models were implemented in Keras, in which the code for the simple model is very concise:

```
model = Sequential()
model.add(Embedding(vocab_size, embed_size,
    input_length=max_title_length))
model.add(Flatten())
model.add(Dropout(0.5))
model.add(Dense(32, activation='softmax'))
model.compile(
    loss='categorical_crossentropy',
    optimizer=Adam())
model.fit(
    trX, to_categorical(trY), epochs=
    num_epochs)
```

As training the models is computationally intensive, we ran many of our tests on Google Colaboratory.

#### V. Experiment

Eleven distinct models were tested. Each model was tuned in terms of hyperparameters and dropout layers through experimentation, so they each represent a fair example of the architecture.

All models were run for at least 25 epochs, with the Adam optimizer at a learning rate of 0.001. We chose a batch size of 128, and an embedding length of 32 (except in the case of using the 100-dimensional GloVe embeddings). Of the 207572 labelled data, we took 9/10 for training and 1/10 for validation.

The results of the naive model, in which no meaning is embedded in the input vector, were predictably poor. The MLP achieved only 9% validation accuracy, while the CNN eventually reached 12%. While these accuracies are better than chance ( $\approx 3\%$ ), they are far from useful. There needs to be a meaningful representation of the input words in the input vectors, not arbitrary word indices.

Adding an embedding layer drastically improved the results. The best 7 models performed very similarly, all falling between 62 and 66%.

After training the model, we tested the dataset by category to get a prediction vector which represented the mean of the predictions for all titles in the validation set belonging to that category. Projecting these 32-dimensional means into




Figure 2. Categorization accuracy for 7 best models

2 dimensions using T-SNE, we can see that meaningful spatial relationships between categories have emerged. For instance, fiction categories Literature, Romance, and Mystery are close to each other, but far away from personal growth categories like Parenting, Self-Help, and Fitness. These relationships can also be seen in the top 5 categorizations for titles in a given category (see Appendix).

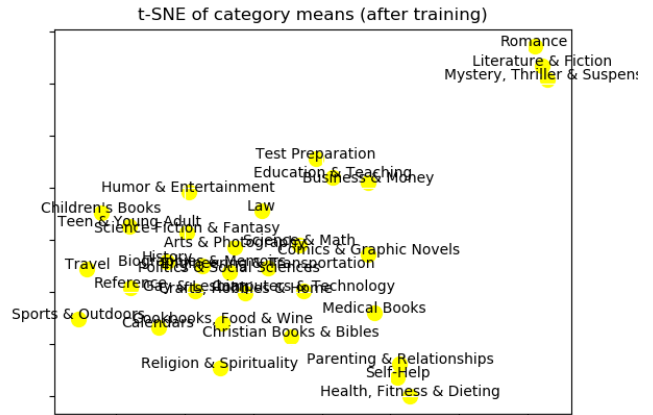


Figure 3. Spatial relationships between categories

#### Training embeddings from dataset

The simple model (code shown above) simply trained the embedding for the vocabulary of the book titles, creating relationships between individual word embeddings and the book categories in the process. The input word vectors, previously organized in rows (one vector per word) were then flattened into a single input vector for the title and fully connected to the 32 category softmax output. This worked

very well ( 80 to 90% accuracy) in certain categories for which specific words consistently appeared, like Calendars.

Adding fully connected layers to create a multi-layer perceptron did not improve this model. The MLP variant was prone to overfitting, and was slower to propagate error correction through to the embedding layer.

The CNN model also didn't outperform the simple model. This suggests that the embedding layer, already taking advantage of word context and the labels of the specific dataset, is sufficient for the task.

### Using pre-trained GloVe embeddings

The GloVe embeddings contain a vocabulary of 400,000 words. Since we disallowed training for this model, that meant words not found in the dataset would not be placed meaningfully in the embedding space. The missing words were set to the 0 vector. We did not explore whether the disclosure of rare words is a negative or positive for this problem, but this is the subject of other researchs in NLP. Out of the 75,094 words in the titles, 20,429 (27%) were not found in GloVe dataset.

Performing experiments on the pre-trained GloVe embeddings better supported our hypothesis that a CNN would be a good choice for this problem. With the simple model (title matrices flattened and fully connected to the softmax output), there was insufficient trainable parameters (83,232) to fit the data to the labels, reaching a plateau at roughly 50% accuracy on the validation data, and less than 60% even on the training data.




Figure 4. GloVe embeddings straight into softmax

The MLP model did little to improve validation accuracy, reaching 54% after 14 epochs. While the millions of

trainable parameters (3,204,640) improved training accuracy to 80% at 50 epochs, the patterns it was learning were not generalizable.




Figure 5. GloVe embeddings through MLP

Adding the convolutional model suggested by Britz (2015a) improved the validation accuracy to the point of being competitive with the embeddings trained from the titles (63%) while having a modest number of trainable parameters (433,184). This was the most convincing result in support of our hypothesis that the CNN would improve categorization accuracy. In contrast with the MLP above, the training clearly had better ability to generalize to unseen titles. This is not surprising, as the MLP would be considering the position of words in the title, whereas the CNN would be looking for words, or combinations of words, in any location in the title, disregarding precisely where in the title they appear. This property of spatial invariance is crucial for creating closer associations between individual word vectors and the category labels.

### Re-training GloVe embeddings from dataset

In this model, we initialized our embeddings to the pre-trained GloVe set and then re-trained based on the titles and labels. Instead of initializing the missing words to zero, we initialized them randomly and allowed them to be trained along with the rest of the pre-trained embeddings. The three tested variations on this model performed the best of all tested models.

The simple model converged to a peak validation accuracy of 64.7% in 9 epochs. The training accuracy continues to increase at 50 epochs, showing similarity to the MLP with non-trainable embeddings discussed previously. This makes sense, as this model also has a huge number of trainable

parameters (7,592,632).



Figure 6. GloVe initialized but trainable

Adding fully connected layers didn't much improve on this result, reaching around the same accuracy (64.8%) but taking 24 epochs to do so. As previously discussed, there are sufficient trainable parameters in the embedding to categorize the data - the addition of more parameters only slows down learning.

The CNN model performed the best of all tested models, with 66% validation accuracy at 20 epochs, and the best accuracy at any epoch of any model. Since the filters can learn spatially invariant word associations with categories, the relationship between words and these categories does not need to be captured in the word embedding. Thus, the generalized semantic relationships of the words, captured in the GloVe training on a much larger dataset, is retained. Apart from the slow training on 7,942,584 parameters, this model is the most promising.

### Factors in the dataset affecting categorization accuracy

Aside from the particular configuration of the models, certain factors in the dataset affected the accuracy. The most substantial were found to be title length and category. Titles of less than 5 words had noticeably worse accuracy, with one word titles being categorized correctly only 30% of the time. However, accuracy did not substantially improve as the length increased over 5 words. (The figure was generated from the simple embedding model).

The accuracy also varied by category. Based on manual inspection of the data, certain categories have distinct words which consistently appear in the titles (like "Calendar" or "Cookbook") which directly mirror the category. Other

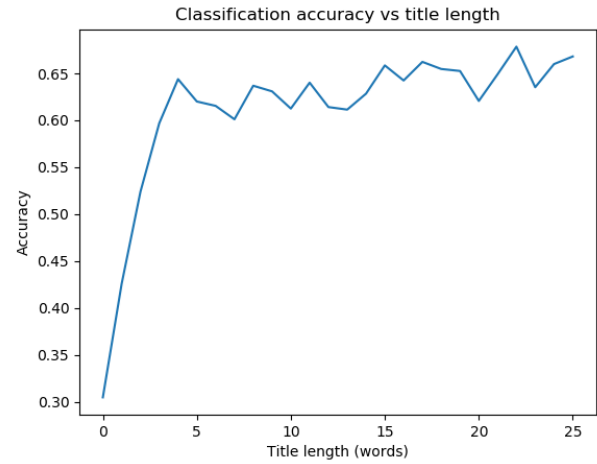


Figure 7. Classification accuracy is a function of length

highly accurate categories like Computers, Medical, and Law often include domain-specific words in the titles. The worst performing categories (Biographies, Parenting, Politics, Self-Help) have more varied titles.




Figure 8. Certain categories achieve better accuracy

Titles from certain categories also exhibit more overlap in terms of potential categorization. By looking at the T-SNE plot of category means, and the accompanying table of categorical output by categorical input, we can see that certain categories cluster closely together and have high error between them. The top 5 accuracy for these categories is still fairly good.

Finally, such a system could never achieve 100% accuracy while there are errors or inconsistency in the training and test data. The title "Carrying off the Cherokee: History of Buff-

ington's Company Georgia Mounted Militia", which the network categorized as History, was actually labelled "Crafts, Hobbies & Home". Other books do not fit cleanly into one category or another, such as "Indonesia (Country Explorers)", which the network categorized as Travel. The given label was "Children's Books". This title demonstrates how a children's travel book would never achieve good categorization accuracy, since one could reasonably place it in either category.

## VI. Conclusion

In this project, we experimented using different neural network models to categorize a book based on its title. We achieved better results using the titles than previous work on categorizing the same titles by their cover images (Iwana et al.). We have shown that with word embedding and a very simple neural network model, we are able to categorize the books fairly accurately, but accuracy is highly dependent on the book category. However, since the neural network is very small, we are able to train the network within a reasonable short time. This makes it possible to train the network with a larger dataset in order to improve the result, especially for those categories which had more diverse titles. Future work could involve a more complicated CNN, or finding a new neural network model that is better at this job, such as a RNN - but a more complex model would make it harder to train the network with a large data set because of the computational complexity. Another possibility to improve the result is to filter out some of the confusing words, or misleading words among the titles, where these words are the most common words that will cause a book title to be incorrectly categorized. NLP research involves experimenting with much pre-processing which fell outside scope of this project, but would be interesting to apply on top of our current model.

The *Book32* dataset is an interesting dataset to be working with, as it provides the capability to conducting many different variations of this classification experiment, like the relationship between a book title and the book author, etc. However, if the original dataset could provide a typical sentence from each book (like the first sentence from the book or the best sentence from the book), then we would expect even better results. Also, it will be helpful if people can enrich that dataset with more books that belong those small sized categories like Education & Training, which will definitely improve the classification accuracy of our model.

Returning to our initial problem statement, we conclude that convolution is indeed a powerful tool for NLP tasks, especially when starting from a pre-trained word embedding. Secondly, we find that initializing a trainable embedding to known meaningful values (such as the GloVe dataset) is a good starting point for an effective NLP neural network.

## References

- (1) Britz, D. (2015a, November 7). Understanding Convolutional Neural Networks for NLP. Retrieved December 11, 2018, from <http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>
- (2) Britz, D. (2015b, December 11). Implementing a CNN for Text Classification in TensorFlow. Retrieved December 11, 2018, from <http://www.wildml.com/2015/12/implementing-a-cnn-for-text-classification-in-tensorflow/>
- (3) GloVe: Global Vectors for Word Representation. (n.d.). Retrieved December 15, 2018, from <https://nlp.stanford.edu/projects/glove/>
- (4) Iwana, B. K., Rizvi, S. T. R., Ahmed, S., Dengel, A., & Uchida, S. (2016). Judging a Book By its Cover. ArXiv:1610.09204 [Cs]. Retrieved from <http://arxiv.org/abs/1610.09204>
- (5) Karani, D. (2018, September 1). Introduction to Word Embedding and Word2Vec. Retrieved December 11, 2018, from <https://towardsdatascience.com/introduction-to-word-embedding-and-word2vec-652d0c2060fa>
- (6) Kim, Y. (2014). Convolutional Neural Networks for Sentence Classification. ArXiv:1408.5882 [Cs]. Retrieved from <http://arxiv.org/abs/1408.5882>
- (7) Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. ArXiv:1301.3781 [Cs]. Retrieved from <http://arxiv.org/abs/1301.3781>
- (8) Peterson, V. (n.d.). How to Write an Effective Book Title. Retrieved December 15, 2018, from <https://www.thebalancecareers.com/how-to-write-a-book-title-insights-and-examples-2799918>
- (9) This dataset contains 207,572 books from the Amazon.com, Inc. marketplace.: uchidalab/book-dataset. (2018). Python, 九州大学 ヒューマンインタフェース研究室. Retrieved from <https://github.com/uchidalab/book-dataset> (Original work published 2017)
- (10) Word2vec. (2018). In Wikipedia. Retrieved from <https://en.wikipedia.org/w/index.php?title=Word2vec&oldid=869116438>
- (11) Zhang, Y., & Wallace, B. (2015). A Sensitivity Analysis of (and Practitioners' Guide to) Convolutional Neural Networks for Sentence Classification. ArXiv:1510.03820 [Cs]. Retrieved from <http://arxiv.org/abs/1510.03820>

- (11) Chollet, F. (2016) Using pre-trained word embeddings in a Keras model. Retrieved December 15, 2018, from <https://blog.keras.io/using-pre-trained-word-embeddings-in-a-keras-model.html>



## Appendix

Additional results from our experiments are shown here. See **V. Experiment** for details and analysis.




Figure 9. Validation accuracy for all 11 tested models




Figure 10. Validation accuracy for 7 best models (detail)




Figure 11. T-SNE projection of means of all validation titles by category




Figure 12. Certain categories achieve better accuracy

3 epochs of simple model, achieving 61% total accuracy

Category	Top 1	Top 5
0. Calendars	0.8937	0.9488
1. Cookbooks, Food & Wine	0.8331	0.9466
2. Travel	0.8199	0.9375
3. Computers & Technology	0.7862	0.9283
4. Medical Books	0.7514	0.9275
5. Test Preparation	0.7466	0.8836
6. Law	0.7125	0.8889
7. Business & Money	0.6933	0.9005
8. Crafts, Hobbies & Home	0.6610	0.8795
9. Children's Books	0.6559	0.9274
10. Health, Fitness & Dieting	0.6545	0.8887
11. Christian Books & Bibles	0.6513	0.8713
12. Science & Math	0.6251	0.8849
13. History	0.6197	0.8825
14. Religion & Spirituality	0.5378	0.8149
15. Comics & Graphic Novels	0.5265	0.7682
16. Sports & Outdoors	0.5250	0.8107
17. Literature & Fiction	0.5039	0.8548
18. Science Fiction & Fantasy	0.4874	0.7479
19. Humor & Entertainment	0.4763	0.8038
20. Romance	0.4736	0.7841
21. Arts & Photography	0.4676	0.7956
22. Engineering & Transportation	0.4205	0.7311
23. Mystery, Thriller & Suspense	0.3865	0.6908
24. Education & Teaching	0.3851	0.7516
25. Teen & Young Adult	0.3630	0.7982
26. Reference	0.3438	0.6341
27. Gay & Lesbian	0.3258	0.5076
28. Biographies & Memoirs	0.3175	0.7438
29. Parenting & Relationships	0.3038	0.6878
30. Politics & Social Sciences	0.2676	0.6704
31. Self-Help	0.2525	0.6777

Science & Math		Computers & Technology	0.7190
-----		Business & Money	0.0545
Science & Math	0.5566	Science & Math	0.0297
Medical Books	0.0481	Medical Books	0.0178
Children's Books	0.0311	Engineering & Transportation	0.0169
Travel	0.0272		
Business & Money	0.0268	Parenting & Relationships	
		-----	
Engineering & Transportation		Parenting & Relationships	0.2721
-----		Health, Fitness & Dieting	0.1171
Engineering & Transportation	0.3116	Self-Help	0.0714
Science & Math	0.1271	Christian Books & Bibles	0.0669
Business & Money	0.0794	Medical Books	0.0623
Computers & Technology	0.0632		
Crafts, Hobbies & Home	0.0550	Religion & Spirituality	
		-----	
Christian Books & Bibles		Religion & Spirituality	0.4682
-----		Christian Books & Bibles	0.1161
Christian Books & Bibles	0.5487	History	0.0347
Religion & Spirituality	0.0976	Literature & Fiction	0.0333
Literature & Fiction	0.0287	Politics & Social Sciences	0.0301
Children's Books	0.0219		
Romance	0.0215	Self-Help	
		-----	
Travel		Self-Help	0.2236
-----		Health, Fitness & Dieting	0.1350
Travel	0.7565	Religion & Spirituality	0.0714
History	0.0353	Medical Books	0.0630
Sports & Outdoors	0.0226	Parenting & Relationships	0.0598
Children's Books	0.0170		
Literature & Fiction	0.0133	Children's Books	
		-----	
Literature & Fiction		Children's Books	0.5142
-----		Teen & Young Adult	0.0566
Literature & Fiction	0.3369	Literature & Fiction	0.0376
Children's Books	0.0603	Humor & Entertainment	0.0329
Romance	0.0583	Sports & Outdoors	0.0291
Teen & Young Adult	0.0511		
Science Fiction & Fantasy	0.0484	Biographies & Memoirs	
		-----	
Sports & Outdoors		Biographies & Memoirs	0.1805
-----		History	0.0981
Sports & Outdoors	0.4798	Children's Books	0.0739
Travel	0.0864	Teen & Young Adult	0.0650
Children's Books	0.0579	Literature & Fiction	0.0596
Biographies & Memoirs	0.0359		
Health, Fitness & Dieting	0.0332	Reference	
		-----	
Computers & Technology		Reference	0.3131
-----		Travel	0.0726
		Children's Books	0.0620
		Humor & Entertainment	0.0516

Teen & Young Adult	0.0347
Cookbooks, Food & Wine	
-----	
Cookbooks, Food & Wine	0.7309
Health, Fitness & Dieting	0.0694
Travel	0.0237
Crafts, Hobbies & Home	0.0217
Children's Books	0.0190
Arts & Photography	
-----	
Arts & Photography	0.3769
Crafts, Hobbies & Home	0.0852
Humor & Entertainment	0.0536
Travel	0.0372
Children's Books	0.0357
Education & Teaching	
-----	
Education & Teaching	0.3207
Business & Money	0.1194
Test Preparation	0.0719
Law	0.0718
Medical Books	0.0532
Law	
-----	
Law	0.6486
Business & Money	0.0818
Medical Books	0.0346
Politics & Social Sciences	0.0240
History	0.0166
Comics & Graphic Novels	
-----	
Comics & Graphic Novels	0.4597
Children's Books	0.0852
Literature & Fiction	0.0475
Teen & Young Adult	0.0468
Science Fiction & Fantasy	0.0444
Science Fiction & Fantasy	
-----	
Science Fiction & Fantasy	0.4511
Literature & Fiction	0.0844
Romance	0.0596
Children's Books	0.0575
Teen & Young Adult	0.0523
Medical Books	

-----	
Medical Books	0.6854
Health, Fitness & Dieting	0.0731
Business & Money	0.0321
Science & Math	0.0316
Politics & Social Sciences	0.0182
Health, Fitness & Dieting	
-----	
Health, Fitness & Dieting	0.6021
Medical Books	0.0766
Cookbooks, Food & Wine	0.0447
Self-Help	0.0289
Teen & Young Adult	0.0240
Gay & Lesbian	
-----	
Gay & Lesbian	0.2466
Literature & Fiction	0.0958
Biographies & Memoirs	0.0655
Politics & Social Sciences	0.0578
Teen & Young Adult	0.0487
History	
-----	
History	0.4471
Travel	0.0944
Biographies & Memoirs	0.0469
Religion & Spirituality	0.0396
Literature & Fiction	0.0345
Calendars	
-----	
Calendars	0.8915
Crafts, Hobbies & Home	0.0217
Travel	0.0127
Humor & Entertainment	0.0089
Arts & Photography	0.0083
Mystery, Thriller & Suspense	
-----	
Mystery, Thriller & Suspense	0.2656
Literature & Fiction	0.1817
Romance	0.0870
Science Fiction & Fantasy	0.0643
Children's Books	0.0568
Politics & Social Sciences	
-----	
Politics & Social Sciences	0.1601
Medical Books	0.0706
Religion & Spirituality	0.0619

Science & Math	0.0616
Business & Money	0.0596
-----	
Business & Money	
-----	
Business & Money	0.5704
Law	0.0487
Computers & Technology	0.0428
Medical Books	0.0354
Science & Math	0.0259
Test Preparation	
-----	
Test Preparation	0.6950
Medical Books	0.0547
Education & Teaching	0.0478
Business & Money	0.0398
Law	0.0336
Humor & Entertainment	
-----	
Humor & Entertainment	0.3760
Children's Books	0.0617
Literature & Fiction	0.0428
Arts & Photography	0.0419
Crafts, Hobbies & Home	0.0339

# Teen & Young Adult

Teen & Young Adult	0.2998
Children's Books	0.1086
Literature & Fiction	0.0640
Health, Fitness & Dieting	0.0451
Biographies & Memoirs	0.0307

# Crafts, Hobbies & Home

Crafts, Hobbies & Home	0.6123
Arts & Photography	0.0592
Science & Math	0.0308
Humor & Entertainment	0.0294
Children's Books	0.0290

# Romance

Romance	0.3775
Literature & Fiction	0.1140
Science Fiction & Fantasy	0.0624
Teen & Young Adult	0.0504
Children's Books	0.0451